

Introduction/Tutorial on the Linux Ecosystem

Hyperion cluster, SLURM scheduler,
Measuring execution time (part 3)

Alexandru Iulian Orhean
aorhean@hawk.iit.edu

Illinois Institute of Technology
Computer Science Department
Data-Intensive Distributed Systems Laboratory



Table of contents

1. Hyperion cluster
2. SLURM scheduler
3. Measuring execution time
4. Hyperion upgrade

Hyperion cluster

Hyperion cluster

- The platform where you will run PA1;
- The platform where we will grade PA1;

```
localhost$ ssh <username>@129.114.33.105  
hyperion$ passwd
```

- SLURM - free & open-source job scheduler for Linux and UNIX*;
- allocates computer resources to users in an **exclusive** and/or non-exclusive mode, for a limited amount of time;
- provides a framework for starting, executing and monitoring work on a set of allocated nodes;
- arbitrates contention for resources by managing a job queue;

SLURM scheduler

SLURM commands

```
$ sinfo
```

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
interactive	up	1:00:00	10	idle	bluecompute-[1-10]
compute*	up	15:00	50	idle	redcompute-[1-50]

```
$ squeue
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST
46	compute	bash	aorhean	R	0:40	1	redcompute-1

```
$ scancel 46
```

```
$ srun -n 1 -p interactive --pty /bin/bash
```

```
$ sbatch run.slurm
```

```
#!/bin/bash

#SBATCH --nodes=2
#SBATCH --output=main.out
#SBATCH --wait-all-nodes=1

echo $SLURM_JOB_NODELIST

./myprogram /tmp/input.txt /tmp/output.txt
```

Measuring execution time

User time vs System time vs Wall time

```
$ time du -sh /home  
8.4G    /home/
```

```
real    0m17.233s  
user    0m0.350s  
sys     0m1.850s
```

- Wall (Real) time -> total time from start to finish, including wait time (end of process quanta or waiting for I/O to complete);
- User time -> total time spent on the CPU in user space (other processes and time the processes spends blocked do not count);
- Sys time -> total time spent on the CPU in kernel space (other processes and time the processes spends blocked do not count);

C example - CPU time

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <unistd.h>

int main(int argc, char **argv)
{
    clock_t start, end;
    start = clock();
    sleep(10);
    end = clock();

    printf("elapsed: %fs\n", (((float) end - start)
        / CLOCKS_PER_SEC));
    return 0;
}
```

```
$ gcc -Wall main.c
$ ./a.out
elapsed: 0.000034s
```

C example - high precision Wall time

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/time.h>
#include <unistd.h>

int main(int argc, char **argv)
{
    struct timeval start, end;
    gettimeofday(&start, NULL);
    sleep(10);
    gettimeofday(&end, NULL);

    printf("elapsed: %fs\n",
           (float) (end.tv_usec - start.tv_usec) / 1000000 +
           (float) (end.tv_sec - start.tv_sec));
    return 0;
}

$ gcc -Wall main.c
$ ./a.out
elapsed: 10.000120s
```

Hyperion upgrade

Friday 2nd of Feb scheduled upgrade!

- not accessible between 6pm-12am;
- upgrade the login node
(more ram & more cores)
- add more compute nodes
(aiming for a total of 90 compute nodes)